Deploying Apache Airflow on an EC2 Instance

JUNE 2024

ISSUED BY
Medha Prodduturi

Introduction

This guide will walk you through the process of connecting AWS EC2 Instance with Visual Studio Code followed by setting up Apache Airflow on an AWS EC2 Instance, allowing you to effectively orchestrate complex computation workflows and data processing pipelines.

In this article:

Connecting VSCode with EC2

Prerequisites

1. An active AWS account.
2. Visual Studio Code installed on your local machine.
3. An EC2 instance launched and running within your AWS account.

4. A downloaded .pem key-pair file following EC2 launch
5. Basic understanding of Apache Airflow, AWS EC2 Instance, and Visual Studio Code.

Step 1: Download Remote-SSH

Navigate to Extensions (left-side) in VSCode and download the extension "Remote-SSH' by Microsoft (shown below).

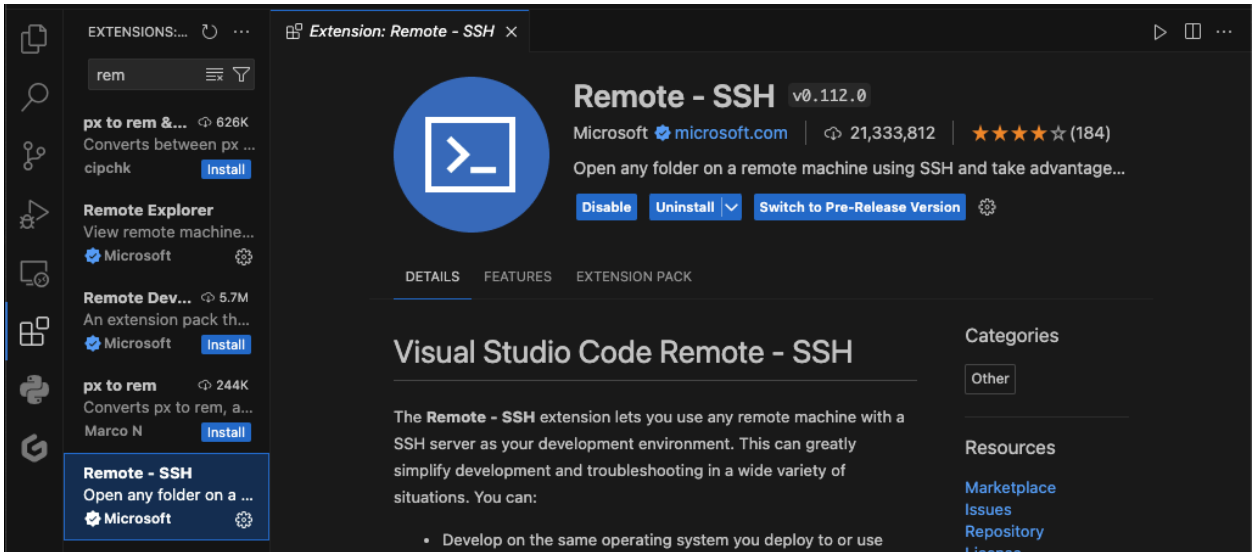

Step 2: Open a Remote Window

Click the blue >< symbols located on the bottom left corner of VSCode to open a remote window.

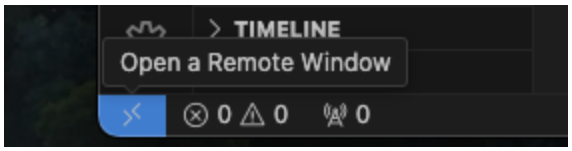

Step 3: Connect & Configure Host

- Click on "Connect to Host"

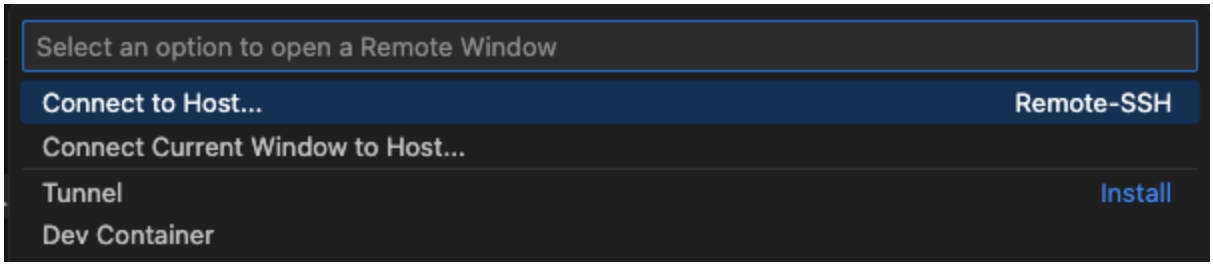

- Click on “Configure SSH Hosts”

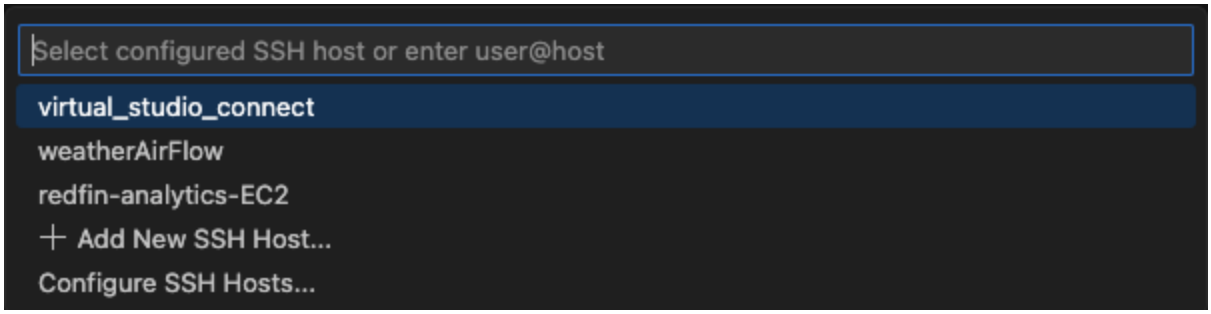

- Select the config folder

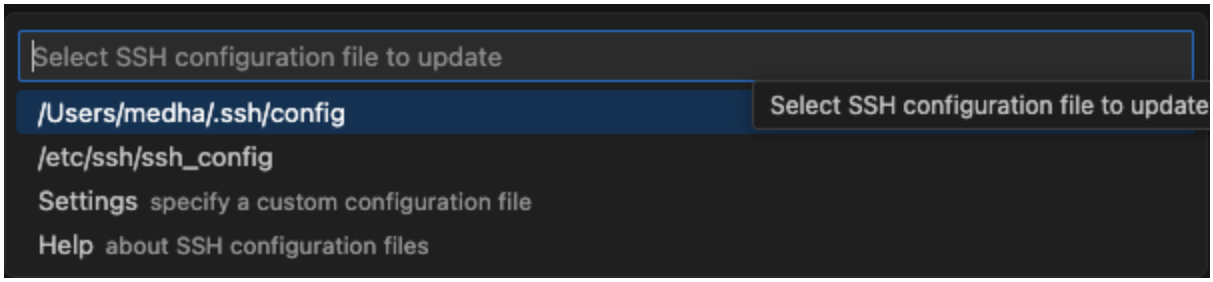

Step 4: Input EC2 Instance details

In the config file, enter the following code according to your EC2 Instance & save it.

```
#Replace this with your EC2 Instance name
Host weatherAirFlow
   # Replace this hostname with your EC2 Public IP address
   HostName 00.000.000.000
   User ubuntu
   #Replace this with your pem file location
   IdentityFile /Users/medha/Downloads/redfin-keypair.pem
```

Step 5: Connect to EC2 Instance

Finally, to connect your EC2 Instance, click the blue >< symbols again.

- Click on “Connect to Host” again.

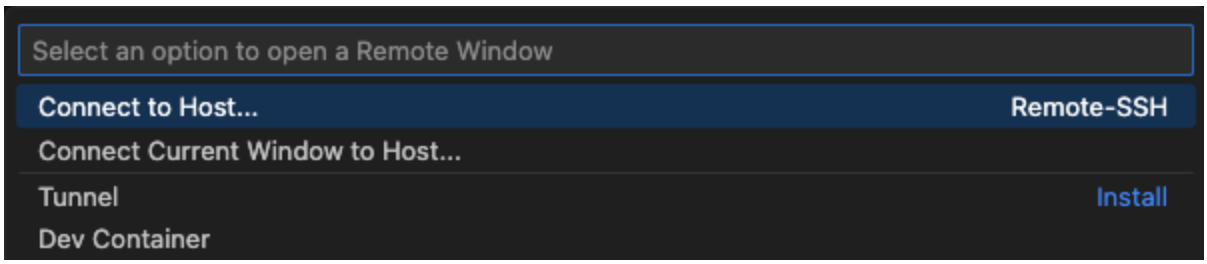

- Choose the designated Host mentioned in your config file (Step 4). In this case, it is weatherAirFlow.

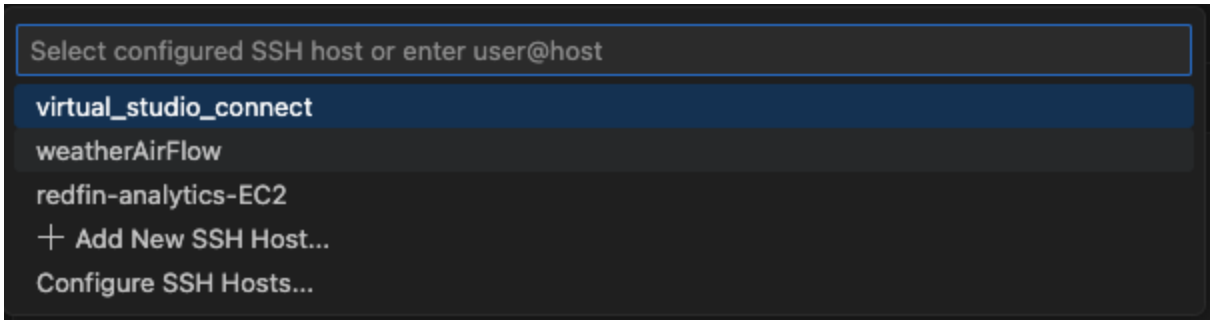

- Successful connection should look like this:

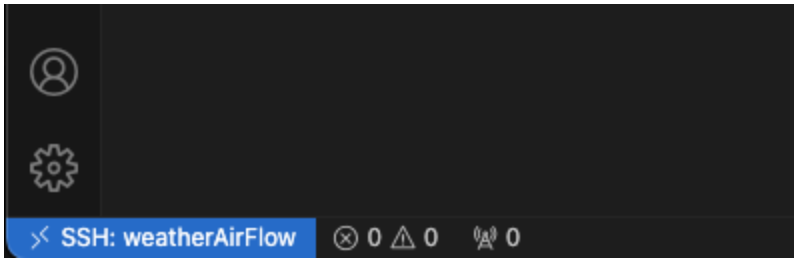

Setting Up Apache Airflow

Prerequisites

1. An active AWS account.
2. Visual Studio Code installed on your local machine.
3. An EC2 Instance launched and running within your AWS account.

Step 1: Connect your EC2 Instance

- Above the list of your instances, you will see the below navigation bar. Click on "Connect"

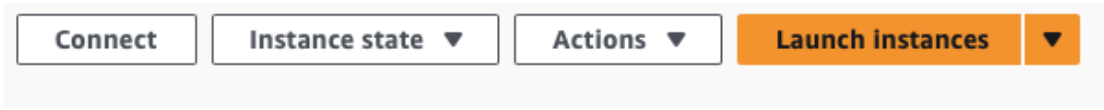

- Clicking on “Connect” will direct you to the below page. Click on “Connect” again.

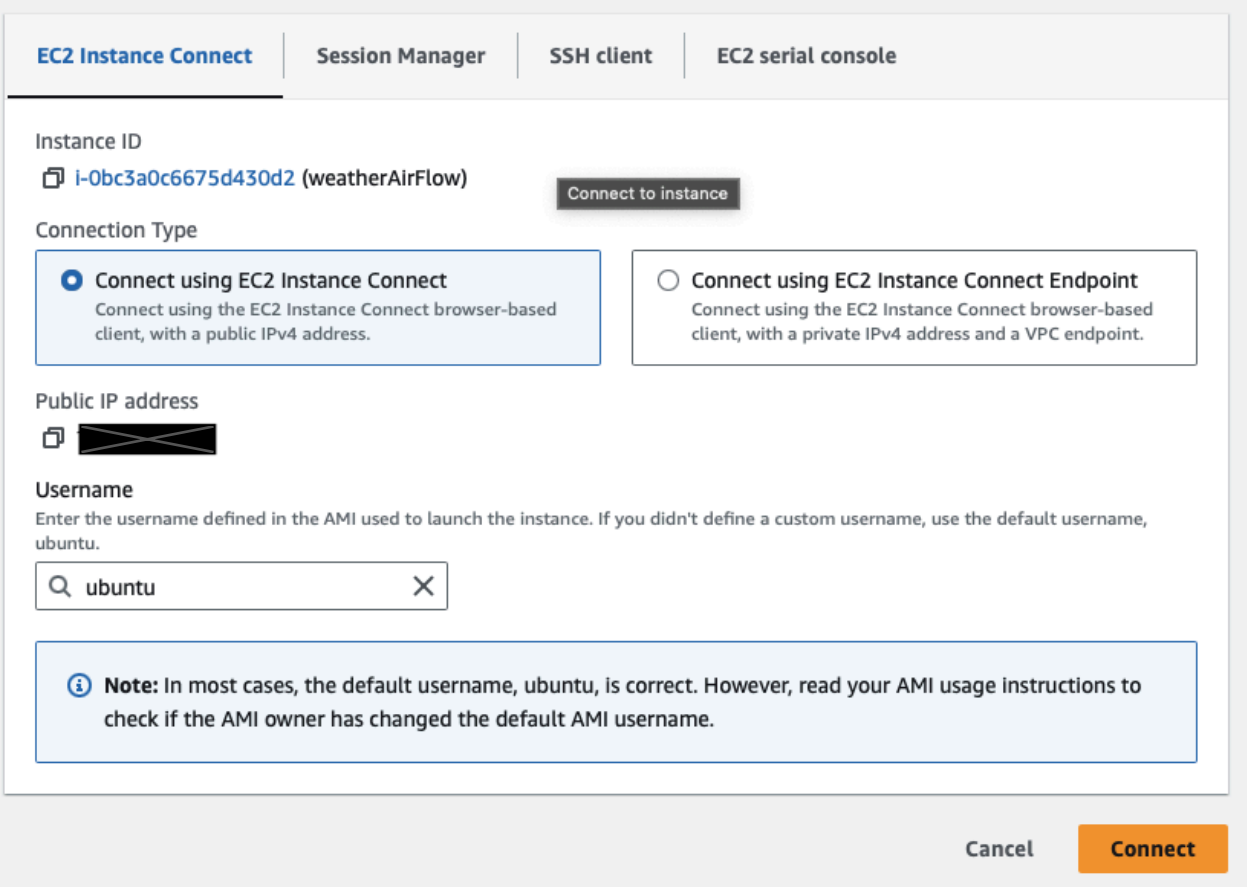

Step 2: Update the package list

Run this command to update existing packages: `sudo apt-get update`

Step 3: Install Python3 package manager

Run this command to install python3 package manager: `sudo apt install python3-pip`

Step 4: Install Apache Airflow

Run this command to install airflow: `sudo pip install apache-airflow`

Step 5: Install Apache Airflow

Install Amazon provider for Apache Airflow: `sudo pip install apache-airflow-providers-amazon`

Step 6: Install Apache in a standalone mode

Run this command: `airflow standalone`

Step 7: Navigate to the 8080 port

In a new tab, enter "[Your EC2 Public IP Address]:8080." This should launch a page like below:

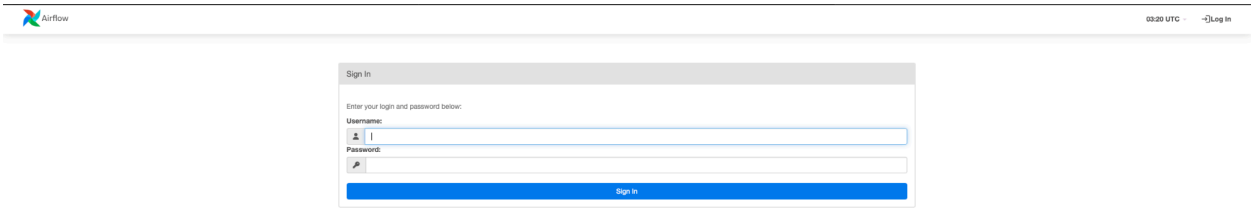

You will find your login credentials at the end of the terminal. Enter those into the page above, and you're in!

Conclusion

By following this guide, you have laid a solid foundation for using Apache Airflow to orchestrate your workflows on AWS. Continue exploring and experimenting to fully leverage the capabilities of Airflow in your data engineering and automation projects.